

## Homework 6

Group 45

Prachi Patel

Jignasuben Vekariya

424-394-6096 (Tel. of Prachi)

857-262-0803 (Tel of Jignasu)

[patel.prachi2@northeastern.edu](mailto:patel.prachi2@northeastern.edu)

[vekariya.j@northeastern.edu](mailto:vekariya.j@northeastern.edu)

Percentage of Effort Contributed by Student 1: 50%

Percentage of Effort Contributed by Student 2: 50%

Signature of Student 1: *Prachi Patel*

Signature of Student 2: *J.R.Vekariya*

Submission Date: 04/27/2022

# Homework\_6 (1) (1)

April 27, 2022

## 1 Homework 6

**Before you start:** Read Chapter 10 Logistic Regression and Chapter 11 Neural Networks in the textbook.

**Note:** Please enter the code along with your comments in the **TODO** section.

Alternative solutions are always welcomed.

```
[ ]: # # Please remove # and run the following code if you have an error while  
→importing the dataset  
# !pip install --upgrade openpyxl
```

### 1.1 Part 1: Logistic Regression

#### 1.1.1 Problem 1 - Financial Condition of Banks

The file **Banks.csv** includes data on a sample of 20 banks.

The “Financial Condition” column records the judgment of an expert on the financial condition of each bank. This response variable takes one of two possible values—weak or strong—according to the financial condition of the bank.

The predictors are two ratios used in the financial analysis of banks: TotLns&Lses/Assets is the ratio of total loans and leases to total assets and TotExp/Assets is the ratio of total expenses to total assets.

The target is to classify the financial condition of a new bank using the two ratios.

```
[126]: import plotly.graph_objs as go  
import plotly.express as px  
from plotly.subplots import make_subplots  
  
import pandas as pd # Loading the required packages required for  
→analysis  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
import math  
import random
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import Ridge

from collections import Counter
from sklearn.neural_network import MLPRegressor

import scipy.stats as ss
import sklearn.preprocessing as sp

from sklearn.multiclass import OneVsOneClassifier

from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler

from sklearn import svm
from sklearn.svm import SVC

from sklearn.ensemble import RandomForestClassifier
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.pipeline import Pipeline
from sklearn import metrics
from sklearn.metrics import classification_report, confusion_matrix, \
    ↳accuracy_score, r2_score, mean_squared_error

from sklearn.model_selection import train_test_split
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer

#\
#!pip3 install catboost
import warnings
warnings.filterwarnings('ignore')
sns.set_style('darkgrid')
cmap = sns.cm.mako_r
%matplotlib inline

from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import \
    ↳BaggingRegressor, AdaBoostRegressor, GradientBoostingRegressor
from sklearn.naive_bayes import MultinomialNB
from sklearn.decomposition import PCA

```

```

from sklearn.pipeline import make_pipeline
from sklearn.datasets import load_iris

from six.moves import input

!pip install mord
!pip install dmbs
from dmbs import classificationSummary, gainsChart, liftChart
from dmbs.metric import AIC_score
import tensorflow as tf

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from keras.models import Sequential
from keras.layers import Dense

```

Requirement already satisfied: mord in /usr/local/lib/python3.7/dist-packages (0.6)

Requirement already satisfied: dmbs in /usr/local/lib/python3.7/dist-packages (0.1.0)

```

[18]: # Import the dataset
from google.colab import files
file = files.upload()
df_bank = pd.read_csv("Banks.csv")
df_bank.head()

```

<IPython.core.display.HTML object>

Saving Banks.csv to Banks (1).csv

```

[18]:
Obs  Financial Condition  TotCap/Assets  TotExp/Assets  TotLns&Lses/Assets
0    1                   1           9.7           0.12           0.65
1    2                   1           1.0           0.11           0.62
2    3                   1           6.9           0.09           1.02
3    4                   1           5.8           0.10           0.67
4    5                   1           4.3           0.11           0.69

```

## TODO 1

Run a logistic regression model (on the entire dataset) that models the status of a bank as a function of the two financial measures provided.

Specify the success class as weak (this is similar to creating a dummy that is 1 for financially weak banks and 0 otherwise), and use the default cutoff value of 0.5.

Let's assume, Weak financial condition=1, strong financial condition=0

```
[127]: y = df_bank['Financial Condition']
x = df_bank.drop(columns=['Obs', 'Financial Condition', 'TotCap/Assets'])
log_reg = LogisticRegression(penalty="l2", C=1e42, solver='liblinear')
log_reg.fit(x,y)
print('Intercept ', log_reg.intercept_[0])
coef=pd.DataFrame({'Coefficient': log_reg.coef_[0]}, index=x.columns)

print(pd.DataFrame({'Coefficient': log_reg.coef_[0]}, index=x.columns))
```

```
Intercept    -14.720832806179043
Coefficient
TotExp/Assets    89.832567
TotLns&Lses/Assets    8.371267
```

## TODO 2

Write the estimated equation that associates the financial condition of a bank with its two predictors in three formats:

- The logit as a function of the predictors
- The odds as a function of the predictors
- The probability as a function of the predictors

$x_1 = \text{TotExp/Assets}$ ,  $x_2 = \text{TotLns\&Lses/Assets}$ ,  $e=2.71828$

A.  $\text{logit} = -14.72083281 + (89.83256659 * x_1) + (8.37126718 * x_2)$

B.  $\text{odds(weak)} = e^{(-14.72083281 + (89.83256659 * x_1) + (8.37126718 * x_2))}$

C.  $p(\text{weak}) = 1 / (1 + (e^{-1(-14.72083281 + (89.83256659 * x_1) + (8.37126718 * x_2))}))$

## TODO 3

Consider a new bank whose total loans and leases/assets ratio = 0.6 and total expenses/assets ratio = 0.11.

From your logistic regression model, estimate the following four quantities for this bank:

the logit, the odds, the probability of being financially weak, and the classification of the bank (use cutoff = 0.5).

```
[128]: logit=log_reg.intercept_[0]+(coef.Coefficient[0]*0.11)+(coef.Coefficient[1]*0.6)
print("logit: ", logit)
odds= math.exp(logit)
print("odds: ", odds)
probability=odds/(1+odds)
print("probability: ", probability)
if (probability<0.5):
    print(" Bank is Financially Strong")
else:
    print("Bank is Financially weak")
```

```
logit: 0.18350982427661666
odds: 1.2014267685026223
probability: 0.5457491412806863
Bank is Financially weak
```

#### TODO 4

We use a cutoff value of 0.5 to classify a record based on propensity.

Instead, if we want to classify the record using the odds or logit, what value should we take as a cutoff?

```
[129]: cutoffs = [0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9]
new_pred = []
for i in cutoffs:
    new_pred=[]
    predict = (log_reg.predict_proba(x) >= i).astype(int)
    for k in predict:
        if k[0] == 1:
            new_pred.append(0)
        else:
            new_pred.append(1)
    print("Accuracy for Cutoff {}: {}".format(i,accuracy_score(y, new_pred)))
```

Accuracy for Cutoff 0.9: 0.85

#### TODO 5

When a bank with in poor financial condition is misclassified as financially strong, the misclassification cost is much higher than a financially strong bank misclassified as weak.

To minimize the expected cost of misclassification, should the cutoff value for classification (which is currently at 0.5) be increased or decreased?

In order to minimise the expected cost of misclassification the cutoff value for the classification should be decreased.A bank that is financially strong may be considered financially weak,which will not affect the bank much,however a financially weak bank has more to loose if it is assumed to be financially strong.Hence reducing the cutoff will reduce the expected cost of misclassification.

### 1.1.2 Problem 2 - Identifying Good System Administrators

A management consultant is studying the roles played by experience and training in a system administrator's ability to complete a set of tasks in a specified amount of time. In particular, the consultant is interested in discriminating between administrators who are able to complete given tasks within a specified time and those who are not.

Data are collected on the performance of 75 randomly selected administrators. They are stored in the file **SystemAdministrators.csv**.

The variable Experience measures months of full-time system administrator experience, while Training measures the number of relevant training credits. The outcome variable Completed is either Yes or No, according to whether or not the administrator completed the tasks.

```
[6]: # Import the dataset
from google.colab import files
file = files.upload()
df1 = pd.read_csv("SystemAdministrators.csv")
df1.head()
```

<IPython.core.display.HTML object>

Saving SystemAdministrators.csv to SystemAdministrators.csv

```
[6]:
```

	Experience	Training	Completed task
0	10.9	4	Yes
1	9.9	4	Yes
2	10.4	6	Yes
3	13.7	6	Yes
4	9.4	8	Yes

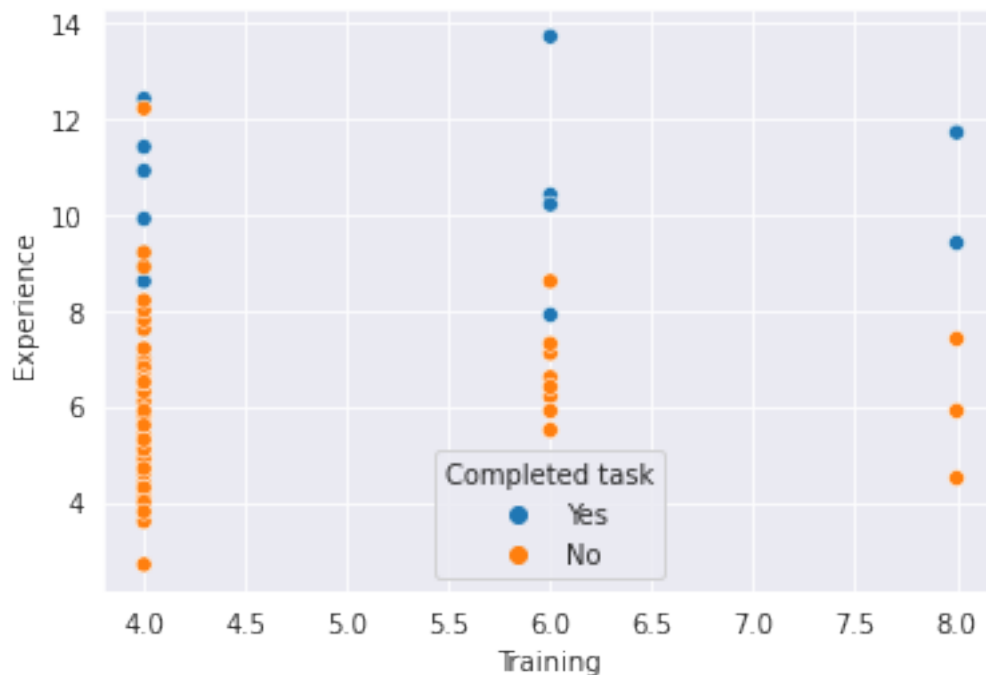
### TODO 1

Create a scatter plot of Experience vs. Training using color or symbol to distinguish the administrators' task completion statuses.

Which predictor(s) appear(s) potentially useful for the classifying task?

```
[130]: Scatter_plot= sns.
↳ scatterplot(x='Training',y='Experience',data=df1,hue='Completed task')
Scatter_plot
```

```
[130]: <matplotlib.axes._subplots.AxesSubplot at 0x7f952902c110>
```



Indented block

## TODO 2

Run a logistic regression model with both predictors using the entire dataset as training data. Among those who completed the task, what is the percentage of administrators incorrectly classified as failing to complete the task?

```
[131]: y=df1['Completed task']
X=df1.drop(['Completed task'], axis=1)
model_log=LogisticRegression()
model_log.fit(X,y)
classes={'No', 'Yes'}
classificationSummary(y, model_log.predict(X), class_names=classes)
```

Confusion Matrix (Accuracy 0.9067)

	Prediction	
Actual	Yes	No
Yes	58	2
No	5	10

```
[132]: cm=confusion_matrix(y, model_log.predict(X))
tn=cm[0,0]
fp=cm[0,1]
tp=cm[1,1]
fn=cm[1,0]
print('The percentage of administators incorrectly classified as failing to_
↪complete the task',cm)
```

The percentage of administators incorrectly classified as failing to complete the task [[58 2]  
[ 5 10]]

58 is indicative of the individuals that correctly completed the task and were correctly predicted while 2 of them correctly completed the task but were incorrectly predicted as failing to complete the task. So, around 2.6 % were incorrectly predicted as failing to complete the task.

## TODO 3

To decrease the percentage in TODO 2, should we increase or decrease the cutoff probability?

The cutoff probability needs to be increased.

## TODO 4

How much experience must be accumulated by an administrator with 4 years of training before his or her estimated probability of completing the task exceeds 0.5?

Since the given training period is 4 years and  $p > 0.5$



```
[133]: p=0.5
odds=p/(1-p)
logit=math.log(odds)
n0=model_log.intercept_[0]
n1=model_log.coef_[0,0]
n2=model_log.coef_[0,1]
exp=(logit-n0-(n2*4))/n1
print('The administrator should have', round(exp,3), 'years experience, so that,
→his or her estimated probability')
```

The administrator should have 9.168 years experience, so that his or her estimated probability

## 1.2 Part 2: Neural Network

### 1.2.1 Problem 3 - Car Sales

Consider the data on used cars (**ToyotaCorolla.csv**) with 1436 records and details on 38 attributes, including Price, Age, KM, HP, and other specifications. The goal is to predict the price of a used Toyota Corolla based on its specifications.

```
[178]: # Import the dataset
from google.colab import files
file = files.upload()
df3 = pd.read_csv("ToyotaCorolla.csv", encoding = 'unicode_escape')
df3.head()
```

<IPython.core.display.HTML object>

Saving ToyotaCorolla.csv to ToyotaCorolla (2).csv

```
[178]:
```

	Id	Model	Price	Age_08_04	\
0	1	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13500	23	
1	2	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13750	23	
2	3	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13950	24	
3	4	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	14950	26	
4	5	TOYOTA Corolla 2.0 D4D HATCHB SOL 2/3-Doors	13750	30	

	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_Color	...	Powered_Windows	\
0	10	2002	46986	Diesel	90	1	...	1	
1	10	2002	72937	Diesel	90	1	...	0	
2	9	2002	41711	Diesel	90	1	...	0	
3	7	2002	48000	Diesel	90	0	...	0	
4	3	2002	38500	Diesel	90	0	...	1	

	Power_Steering	Radio	Mistlamps	Sport_Model	Backseat_Divider	\
0	1	0	0	0	1	
1	1	0	0	0	1	
2	1	0	0	0	1	

3	1	0	0	0	1
4	1	0	1	0	1

	Metallic_Rim	Radio_cassette	Parking_Assistant	Tow_Bar
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 39 columns]

## TODO 1

Fit a neural network model with the following specifications:

- Use predictors Age\_08\_04, KM, Fuel\_Type, HP, Automatic, Doors, Quarterly\_Tax, Mfr\_Guarantee, Guarantee\_Period, Airco, Automatic\_airco, CD\_Player, Powered\_Windows, Sport\_Model, and Tow\_Bar
- Partition the data into 80% training and 20% validation
- Scale the numerical predictor and outcome; convert categorical predictors to dummies
- The neural network should have one single hidden layer with two nodes

Present the model summary and performance on the training and validation sets.

```
[179]: x1 = ['Age_08_04', 'KM', 'Fuel_Type', 'HP', 'Automatic', 'Doors', \
↳ 'Quarterly_Tax', 'Mfr_Guarantee', 'Guarantee_Period', 'Airco',
        'Automatic_airco', 'CD_Player', 'Powered_Windows', 'Sport_Model', \
↳ 'Tow_Bar', 'Price']
x1_data = df3[x1]
x1_data.head()
```

```
[179]:   Age_08_04   KM Fuel_Type  HP  Automatic  Doors  Quarterly_Tax  \
0         23  46986   Diesel  90          0      3           210
1         23  72937   Diesel  90          0      3           210
2         24  41711   Diesel  90          0      3           210
3         26  48000   Diesel  90          0      3           210
4         30  38500   Diesel  90          0      3           210

   Mfr_Guarantee  Guarantee_Period  Airco  Automatic_airco  CD_Player  \
0              0                 3      0              0          0
1              0                 3      1              0          1
2              1                 3      0              0          0
3              1                 3      0              0          0
4              1                 3      1              0          0

   Powered_Windows  Sport_Model  Tow_Bar  Price
0              1          0          0  13500
1              0          0          0  13750
```

2	0	0	0	13950
3	0	0	0	14950
4	1	0	0	13750

```
[180]: x1_data.describe()
```

```
[180]:
```

	Age_08_04	KM	HP	Automatic	Doors \
count	1436.000000	1436.000000	1436.000000	1436.000000	1436.000000
mean	55.947075	68533.259749	101.502089	0.055710	4.033426
std	18.599988	37506.448872	14.981080	0.229441	0.952677
min	1.000000	1.000000	69.000000	0.000000	2.000000
25%	44.000000	43000.000000	90.000000	0.000000	3.000000
50%	61.000000	63389.500000	110.000000	0.000000	4.000000
75%	70.000000	87020.750000	110.000000	0.000000	5.000000
max	80.000000	243000.000000	192.000000	1.000000	5.000000

	Quarterly_Tax	Mfr_Guarantee	Guarantee_Period	Airco \
count	1436.000000	1436.000000	1436.000000	1436.000000
mean	87.122563	0.409471	3.815460	0.508357
std	41.128611	0.491907	3.011025	0.500104
min	19.000000	0.000000	3.000000	0.000000
25%	69.000000	0.000000	3.000000	0.000000
50%	85.000000	0.000000	3.000000	1.000000
75%	85.000000	1.000000	3.000000	1.000000
max	283.000000	1.000000	36.000000	1.000000

	Automatic_airco	CD_Player	Powered_Windows	Sport_Model \
count	1436.000000	1436.000000	1436.000000	1436.000000
mean	0.056407	0.218663	0.561978	0.300139
std	0.230786	0.413483	0.496317	0.458478
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000	0.000000
75%	0.000000	0.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

	Tow_Bar	Price
count	1436.000000	1436.000000
mean	0.277855	10730.824513
std	0.448098	3626.964585
min	0.000000	4350.000000
25%	0.000000	8450.000000
50%	0.000000	9900.000000
75%	1.000000	11950.000000
max	1.000000	32500.000000

```
[181]: x1_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age_08_04             1436 non-null  int64
1   KM                    1436 non-null  int64
2   Fuel_Type             1436 non-null  object
3   HP                    1436 non-null  int64
4   Automatic             1436 non-null  int64
5   Doors                 1436 non-null  int64
6   Quarterly_Tax         1436 non-null  int64
7   Mfr_Guarantee         1436 non-null  int64
8   Guarantee_Period      1436 non-null  int64
9   Airco                 1436 non-null  int64
10  Automatic_airco        1436 non-null  int64
11  CD_Player              1436 non-null  int64
12  Powered_Windows       1436 non-null  int64
13  Sport_Model            1436 non-null  int64
14  Tow_Bar                1436 non-null  int64
15  Price                  1436 non-null  int64
dtypes: int64(15), object(1)
memory usage: 179.6+ KB

```

```

[182]: x1_data.isnull().sum()
x1_data['Fuel_Type'].value_counts()

```

```

[182]: Petrol    1264
Diesel      155
CNG         17
Name: Fuel_Type, dtype: int64

```

```

[183]: f_dum= pd.get_dummies(x1_data['Fuel_Type'])
f_dum.head()

```

```

[183]:
   CNG  Diesel  Petrol
0     0     1     0
1     0     1     0
2     0     1     0
3     0     1     0
4     0     1     0

```

```

[184]: x1_data = pd.concat([x1_data,f_dum], axis=1)
x1_data.head()

```

```

[184]:
   Age_08_04  KM  Fuel_Type  HP  Automatic  Doors  Quarterly_Tax  \
0         23  46986   Diesel  90           0         3           210
1         23  72937   Diesel  90           0         3           210

```

2	24	41711	Diesel	90	0	3	210
3	26	48000	Diesel	90	0	3	210
4	30	38500	Diesel	90	0	3	210

	Mfr_Guarantee	Guarantee_Period	Airco	Automatic_airco	CD_Player	\
0	0		3	0	0	0
1	0		3	1	0	1
2	1		3	0	0	0
3	1		3	0	0	0
4	1		3	1	0	0

	Powered_Windows	Sport_Model	Tow_Bar	Price	CNG	Diesel	Petrol
0	1	0	0	13500	0	1	0
1	0	0	0	13750	0	1	0
2	0	0	0	13950	0	1	0
3	0	0	0	14950	0	1	0
4	1	0	0	13750	0	1	0

```
[185]: x1_data=x1_data.drop(['Fuel_Type'], axis=1)
x1_data.head()
```

```
[185]:
```

	Age_08_04	KM	HP	Automatic	Doors	Quarterly_Tax	Mfr_Guarantee	\
0	23	46986	90	0	3	210	0	
1	23	72937	90	0	3	210	0	
2	24	41711	90	0	3	210	1	
3	26	48000	90	0	3	210	1	
4	30	38500	90	0	3	210	1	

	Guarantee_Period	Airco	Automatic_airco	CD_Player	Powered_Windows	\
0	3	0	0	0	1	
1	3	1	0	1	0	
2	3	0	0	0	0	
3	3	0	0	0	0	
4	3	1	0	0	1	

	Sport_Model	Tow_Bar	Price	CNG	Diesel	Petrol
0	0	0	13500	0	1	0
1	0	0	13750	0	1	0
2	0	0	13950	0	1	0
3	0	0	14950	0	1	0
4	0	0	13750	0	1	0

```
[193]: x1_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype
#   ...
```

```

---  -----
0   Age_08_04      1436 non-null   float64
1   KM             1436 non-null   float64
2   HP             1436 non-null   float64
3   Automatic      1436 non-null   int64
4   Doors          1436 non-null   float64
5   Quarterly_Tax  1436 non-null   float64
6   Mfr_Guarantee  1436 non-null   int64
7   Guarantee_Period 1436 non-null   float64
8   Airco          1436 non-null   int64
9   Automatic_airco 1436 non-null   int64
10  CD_Player      1436 non-null   int64
11  Powered_Windows 1436 non-null   int64
12  Sport_Model    1436 non-null   int64
13  Tow_Bar        1436 non-null   int64
14  Price          1436 non-null   float64
15  CNG            1436 non-null   uint8
16  Diesel         1436 non-null   uint8
17  Petrol         1436 non-null   uint8
dtypes: float64(7), int64(8), uint8(3)
memory usage: 172.6 KB

```

```
[194]: x1_data.isnull().sum()
```

```

[194]: Age_08_04      0
      KM             0
      HP             0
      Automatic      0
      Doors          0
      Quarterly_Tax  0
      Mfr_Guarantee  0
      Guarantee_Period 0
      Airco          0
      Automatic_airco 0
      CD_Player      0
      Powered_Windows 0
      Sport_Model    0
      Tow_Bar        0
      Price          0
      CNG            0
      Diesel         0
      Petrol         0
      dtype: int64

```

```

[195]: num =
      ↪x1_data[['Age_08_04', 'KM', 'HP', 'Doors', 'Quarterly_Tax', 'Guarantee_Period', 'Price']]
      num.head()

```

```
[195]:
```

	Age_08_04	KM	HP	Doors	Quarterly_Tax	Guarantee_Period \
0	-1.771966	-0.574695	-0.768042	-1.085139	2.98868	-0.270919
1	-1.771966	0.117454	-0.768042	-1.085139	2.98868	-0.270919
2	-1.718184	-0.715386	-0.768042	-1.085139	2.98868	-0.270919
3	-1.610620	-0.547650	-0.768042	-1.085139	2.98868	-0.270919
4	-1.395491	-0.801028	-0.768042	-1.085139	2.98868	-0.270919

	Price
0	0.763763
1	0.832715
2	0.887877
3	1.163685
4	0.832715

```
[196]: scaler = StandardScaler().fit(num)
x1_data[['Age_08_04', 'KM', 'HP', 'Doors', 'Quarterly_Tax', 'Guarantee_Period', 'Price']]
↳ scaler.transform(num)
x1_data.head()
```

```
[196]:
```

	Age_08_04	KM	HP	Automatic	Doors	Quarterly_Tax \
0	-1.771966	-0.574695	-0.768042	0	-1.085139	2.98868
1	-1.771966	0.117454	-0.768042	0	-1.085139	2.98868
2	-1.718184	-0.715386	-0.768042	0	-1.085139	2.98868
3	-1.610620	-0.547650	-0.768042	0	-1.085139	2.98868
4	-1.395491	-0.801028	-0.768042	0	-1.085139	2.98868

	Mfr_Guarantee	Guarantee_Period	Airco	Automatic_airco	CD_Player \
0	0	-0.270919	0	0	0
1	0	-0.270919	1	0	1
2	1	-0.270919	0	0	0
3	1	-0.270919	0	0	0
4	1	-0.270919	1	0	0

	Powered_Windows	Sport_Model	Tow_Bar	Price	CNG	Diesel	Petrol
0	1	0	0	0.763763	0	1	0
1	0	0	0	0.832715	0	1	0
2	0	0	0	0.887877	0	1	0
3	0	0	0	1.163685	0	1	0
4	1	0	0	0.832715	0	1	0

```
[197]: b= x1_data['Price']
a = x1_data.drop("Price",axis=1)
a.head()
```

```
[197]:
```

	Age_08_04	KM	HP	Automatic	Doors	Quarterly_Tax \
0	-1.771966	-0.574695	-0.768042	0	-1.085139	2.98868
1	-1.771966	0.117454	-0.768042	0	-1.085139	2.98868

2	-1.718184	-0.715386	-0.768042	0	-1.085139	2.98868
3	-1.610620	-0.547650	-0.768042	0	-1.085139	2.98868
4	-1.395491	-0.801028	-0.768042	0	-1.085139	2.98868

	Mfr_Guarantee	Guarantee_Period	Airco	Automatic_airco	CD_Player	\
0	0	-0.270919	0	0	0	
1	0	-0.270919	1	0	1	
2	1	-0.270919	0	0	0	
3	1	-0.270919	0	0	0	
4	1	-0.270919	1	0	0	

	Powered_Windows	Sport_Model	Tow_Bar	CNG	Diesel	Petrol
0	1	0	0	0	1	0
1	0	0	0	0	1	0
2	0	0	0	0	1	0
3	0	0	0	0	1	0
4	1	0	0	0	1	0

```
[198]: b.head()
```

```
[198]: 0    0.763763
      1    0.832715
      2    0.887877
      3    1.163685
      4    0.832715
      Name: Price, dtype: float64
```

```
[199]: X_train, X_valid, y_train, y_valid = train_test_split(a, b, test_size=0.20,
      ↪ random_state=1)
```

Model 1:

```
[201]: regression = MLPRegressor(hidden_layer_sizes=(2), random_state=1).fit(X_train,
      ↪ y_train)
      regression.predict(X_valid)
```

```
[201]: array([ 4.13137273e-01, -2.78643090e-01,  9.81301075e-01, -1.02931365e+00,
      2.04481649e-01,  2.10965084e+00, -6.56250054e-01, -2.20704720e-01,
      -6.44300638e-01,  1.52723088e-01, -2.17465381e-01, -4.65552513e-01,
      -6.95288710e-01, -6.09565306e-01, -1.11309156e+00, -5.79798635e-01,
      -1.11309156e+00, -8.21241817e-01,  1.82584280e+00, -1.11309156e+00,
      1.27206418e+00, -9.82955762e-01, -5.65679096e-01, -9.84839771e-01,
      4.86708214e-01,  2.73867219e+00, -6.74847923e-01, -2.44774803e-01,
      2.24758923e+00, -3.84419394e-01, -7.91467681e-01, -6.95397648e-01,
      5.43392685e-01,  1.93647388e+00,  4.57893513e+00,  2.95400801e-02,
      -1.80152453e-01, -8.26706043e-01, -5.88177427e-01,  1.80091563e+00,
      7.15184732e-01, -3.83952975e-01, -6.30583922e-01, -6.12443878e-01,
      -6.28500808e-01, -3.77439981e-01,  3.23780163e-01, -4.70126522e-01,
```



-1.11309156e+00, -9.98950485e-01, -6.63469366e-01, 2.87509992e-01,  
 -6.89036619e-01, -9.30942418e-01, -4.95356274e-01, 2.28195811e+00,  
 1.13352614e+00, 3.25799156e-01, -3.11402640e-01, -1.11309156e+00,  
 1.39194953e+00, -4.50820040e-01, -3.72944961e-01, 2.47918207e-01,  
 -9.97782817e-02, 1.21202546e+00, 1.57000466e-01, -3.75160079e-01,  
 -1.62125313e-01, 1.28981056e-01, -1.01652836e+00, -2.00602563e-01,  
 -8.43599199e-01, -7.01383483e-01, -9.65097633e-01, 1.04862752e+00,  
 -6.44889121e-01, -5.64566790e-01, -5.06435444e-01, 5.65033344e-01,  
 4.13966559e-01, -9.00031209e-01, -1.11309156e+00, -2.63606958e-02,  
 -4.00957512e-01, -5.30014427e-01, 3.93839111e-01, 1.51906676e+00,  
 3.74874983e-01, -1.03060873e+00, -3.53594542e-01, 3.76864121e-02,  
 -3.17109226e-01, -2.33828194e-01, -7.07642202e-01, 2.93250507e-01,  
 6.43353049e-01, 7.95571953e-02, 6.52302942e-02, 3.00363411e-02,  
 -6.35241909e-01, -2.16106274e-01, -5.50445630e-02, -7.35771465e-01,  
 2.60290005e-01, 6.86405966e-01, -8.22935320e-02, -7.55798257e-01,  
 -2.74290529e-01, -8.39129724e-02, -8.33112376e-01, 6.06167920e-01,  
 -6.09088343e-01, -1.11309156e+00, -1.80053509e-01, 8.08275632e-01,  
 -1.11309156e+00, -5.74157952e-01, -4.89373395e-01, -1.04943673e+00,  
 2.18523056e+00, 2.58133426e+00, -6.75086794e-01, 2.45512353e-01,  
 -1.00748554e+00, 1.38396321e+00, -1.10412926e+00, 1.02395253e-01,  
 -1.02309999e-01, 2.74620571e+00, 2.37153986e+00, -6.16733478e-01,  
 -3.54631745e-01, -3.60237660e-01, 3.24709991e-01, -4.74478493e-01,  
 -2.27802760e-02, -5.28273310e-01, 8.31269762e-01, -5.76214565e-01,  
 -4.27549642e-01, -4.33133757e-01, -1.78151107e-01, -8.95581243e-01,  
 -1.43955634e-02, 2.58422005e-01, 1.77491952e+00, -2.93372851e-01,  
 -4.22539750e-01, -1.11309156e+00, -7.87991913e-01, -1.02975763e+00,  
 2.40301543e+00, -9.62539923e-01, -1.11309156e+00, 1.40090317e+00,  
 1.25501320e-01, -7.44910958e-01, -7.25166745e-01, 6.94531572e-01,  
 -2.65723140e-01, 4.44022611e-01, -1.11309156e+00, -7.88585436e-01,  
 -8.08260264e-01, -1.11309156e+00, -1.48785706e-01, -3.90557062e-01,  
 -9.08920469e-01, 4.91179466e-01, -3.87958021e-01, 2.77213933e+00,  
 -3.35651117e-01, 1.42605876e-01, 3.41939374e+00, -2.13833811e-01,  
 -7.29425293e-01, -3.02068477e-01, -9.41895454e-01, -8.15630220e-01,  
 6.69539567e-01, 1.83556507e+00, -9.82125934e-01, -9.42514559e-01,  
 -3.61426241e-01, -6.24242446e-01, -5.00539914e-01, -2.61466300e-01,  
 -1.90353386e-01, 1.92292983e+00, 7.67912788e-01, -6.59916617e-01,  
 1.38296247e+00, -2.05388006e-01, -9.55391820e-01, -4.34641096e-01,  
 -5.22162380e-01, -7.85571577e-01, -4.69155458e-01, -8.22204101e-01,  
 -8.67271420e-01, 1.56644393e-02, 2.27260545e+00, -7.78495372e-01,  
 5.01188200e-01, -5.94242835e-01, 4.04691494e-02, -7.20163742e-01,  
 -5.50535254e-01, 1.54817869e+00, -2.08472040e-01, -6.87519881e-01,  
 -3.07884993e-01, -4.77866341e-01, 9.16490634e-01, -9.27010350e-01,  
 -5.67493638e-01, 7.13492203e-01, 5.05911857e-01, -3.31215902e-01,  
 2.28056838e-01, -4.05408963e-01, -6.12724695e-01, -7.50869932e-01,  
 3.83058974e-01, -3.42473250e-01, 1.77993728e+00, -2.22394598e-01,  
 -2.18229913e-01, -1.49142773e-01, 7.64863987e-01, -6.70394217e-01,  
 1.74458872e+00, 5.48889315e-03, -5.28447694e-04, -7.08677035e-01,

```

2.48441287e+00, -6.72499698e-01, 3.95151709e-01, -4.97085850e-01,
-4.38406339e-01, -4.10449214e-01, 1.09662771e+00, 1.18483383e+00,
9.91547021e-02, -3.15543691e-01, -1.99483817e-01, -7.33447447e-01,
-6.87465138e-01, 5.60388592e-01, -6.86232405e-01, 4.80704494e-01,
2.95811167e+00, 5.04192191e-01, -3.39537753e-01, -5.10284023e-01,
-3.47971589e-01, -8.34002867e-01, -7.61888218e-01, 4.37925147e-01,
-2.18337389e-01, 1.55091498e-01, -1.06450993e+00, 2.87535273e+00,
4.54524680e-01, 1.16980506e+00, -3.34500318e-01, -9.86970422e-02,
-1.11309156e+00, -5.22010099e-01, -1.75275564e-01, 2.52492814e+00,
1.04157882e+00, -9.14123378e-02, 2.04324990e+00, -9.51885299e-02,
-7.02860758e-01, -3.72655073e-01, -5.25242810e-01, -5.39762778e-01,
1.73431969e+00, -7.79168268e-01, -9.36091524e-01, -7.87689491e-01,
6.10319380e-01, 1.57616622e+00, -6.60962078e-01, -4.37817118e-01])

```

```
[202]: regression.score(X_train, y_train)
```

```
[202]: 0.8881522743121312
```

```
[203]: regression.score(X_valid, y_valid)
```

```
[203]: 0.9029320302838734
```

High values of  $r^2$  is indicative of a good model

## TODO 2

Repeat the process, changing the number of hidden layers and nodes to {single layer with five nodes}, {two layers, five nodes in each layer}.

Comment on the performance of the three models above.

Model 2 (Single layer with 5 nodes)

```
[204]: r2 = MLPRegressor(hidden_layer_sizes=(5), random_state=1).fit(X_train, y_train)
r2.predict(X_valid)
```

```
[204]: array([ 0.23236521, -0.20808668,  0.89367006, -0.92655803,  0.23168687,
 1.90487523, -0.84932051, -0.50547726, -0.90924789,  0.07565079,
-0.33580599, -0.3105919 , -0.68290915, -0.46946362, -0.81387794,
-0.80068424, -0.87310017, -0.84754929,  1.550843 , -0.66164489,
 1.11037193, -0.64435851, -0.41912993, -0.6764134 ,  0.47814467,
 2.62310348, -0.66361995, -0.15759763,  2.29825755, -0.27975717,
-0.73580543, -0.82693975,  0.5236992 ,  1.71722841,  4.50440188,
-0.12345581, -0.12975393, -0.61015918, -0.44328287,  1.73598824,
 0.59564834, -0.43322167, -0.76986565, -0.60232365, -0.79869015,
-0.25085791,  0.22645084, -0.38355131, -0.82902161, -0.62685907,
-0.86419335,  0.05419903, -0.53556665, -0.62762179, -0.69303409,
 2.13590979,  0.69517411,  0.22807367, -0.35028311, -0.87667689,
 1.41780154, -0.7686446 , -0.67783688,  0.27048428, -0.06568927,
 1.1840564 ,  0.03421106, -0.57950541, -0.17338955, -0.17717222,

```

```

-0.61850424, -0.27670906, -1.05618426, -0.96331634, -0.65568604,
0.9408194 , -0.4533249 , -0.5661388 , -0.77446669, 0.57395884,
0.31494424, -0.98953979, -0.85615348, -0.02442847, -0.55134697,
-0.76186358, 0.87944755, 1.45344391, 0.26140647, -1.25889001,
-0.2303684 , -0.10271356, -0.16205408, -0.3364574 , -0.66410338,
0.1821588 , 0.5502532 , 0.05450607, 0.08341154, -0.04662108,
-0.96948788, -0.19877145, -0.17043431, -0.6764134 , 0.32492171,
0.86605057, -0.08094655, -0.59197899, -0.43676631, -0.2521074 ,
-1.06491662, 0.65013983, -0.9649683 , -0.8329975 , -0.24516994,
0.76094401, -0.71837839, -0.65509403, -0.61299321, -0.61461195,
2.12534096, 3.08747046, -0.55046563, 0.16114465, -0.89060788,
0.95397562, -0.70002139, -0.05147177, -0.06357385, 2.79574261,
2.62972254, -0.75655389, -0.56432865, -0.49580895, 0.35116539,
-0.59554202, -0.03902334, -0.37466684, 0.76061851, -0.6764134 ,
-0.57945878, -0.11840718, -0.29334301, -0.95543555, -0.10696205,
0.22299788, 1.46724886, -0.26895515, -0.15036809, -0.85674724,
-1.00607855, -0.6764134 , 1.88355069, -0.57206374, -0.67438584,
2.00738207, 0.02169663, -0.85899237, -0.86554261, 0.53350567,
-0.37869606, 0.44515499, -0.92124981, -0.67016496, -0.64655589,
-0.46381857, 0.09729428, -0.56355103, -1.0017914 , 0.51802238,
-0.59150672, 2.89530261, -0.2498926 , 0.08355206, 3.75286377,
-0.16003083, -0.60786683, -0.42477446, -0.90854343, -0.76205529,
0.66085814, 1.79636388, -0.62791377, -1.09315039, -0.50094213,
-0.94484588, -0.31995417, -0.3878602 , -0.36286836, 1.64057178,
0.70580482, -0.67247259, 1.48854482, -0.33362603, -0.6311571 ,
-0.37993569, -0.44288884, -0.6764134 , -0.41012794, -0.92084771,
-0.65330502, -0.08651082, 2.23089131, -0.63566996, 0.63547443,
-0.75081997, 0.13838668, -0.82227601, -0.60759293, 1.63205976,
-0.3246709 , -0.64680701, -0.18779366, -0.59122524, 0.78007869,
-0.79921819, -0.31307653, 0.94546666, 0.94952085, -0.32398667,
0.37166833, -0.66055358, -0.54782377, -0.75045621, 0.58180946,
-0.12351008, 1.71232914, -0.42454193, -0.18626812, -0.36361051,
0.71299134, -0.81785878, 1.72325531, -0.01289127, -0.13193451,
-0.81309896, 2.65289603, -0.6764134 , 0.36964809, -0.4379183 ,
-0.70650393, -0.50866415, 1.07238236, 1.10205641, 0.02665728,
-0.22964132, -0.3442441 , -0.60874851, -0.70532612, 0.45256386,
-0.47844625, 0.50213625, 3.13496135, 1.00936622, -0.34128549,
-0.24138141, -0.52596273, -0.82164419, -0.75697468, 0.26855054,
-0.25332322, 0.18628307, -0.77502097, 3.21815616, 0.46330928,
1.21880664, -0.23996349, -0.12594826, -0.95597518, -0.43734385,
-0.0807367 , 2.11160286, 0.98559823, -0.25013762, 2.31005396,
-0.28001896, -0.67452929, -0.36659735, -0.3782435 , -0.89635928,
1.64307342, -0.72957003, -0.89623324, -0.88976469, 0.57570952,
1.62846414, -0.5835138 , -0.75234747])

```

```
[205]: r2.score(X_train, y_train)
```

[205]: 0.8888120738085448

```
[206]: r2.score(X_valid, y_valid)
```

[206]: 0.900462935100548

High values of r2 is indicatice of having a good performing model.

Model 3 (Two Layers with 5 nodes)

```
[207]: r3 = MLPRegressor(hidden_layer_sizes=(5,5), random_state=1).fit(X_train,
    ↪y_train)
r3.predict(X_valid)
```

```
[207]: array([ 2.22405613e-01, -1.70176023e-01,  1.04152492e+00, -7.26240742e-01,
  1.79226735e-01,  2.01510324e+00, -6.87272858e-01, -6.38028716e-01,
 -6.88024616e-01, -5.97698541e-01, -4.57874265e-01, -4.22713102e-01,
 -7.65832794e-01, -7.30644211e-01, -7.65798581e-01, -6.74928480e-01,
 -7.38756084e-01, -6.28824821e-01,  1.63295099e+00, -7.65832794e-01,
  1.20302218e+00, -4.00697653e-01, -4.43484120e-01, -7.65832794e-01,
  5.95640471e-01,  2.52425740e+00, -7.42166485e-01, -3.17632581e-01,
  2.08699116e+00, -8.32950198e-01, -6.91141080e-01, -7.01903589e-01,
  7.03706092e-01,  1.77363285e+00,  4.24891659e+00, -2.92118976e-01,
 -2.94197711e-01, -4.55937287e-01, -5.04145234e-01,  1.69342387e+00,
  5.70883677e-01, -4.82749123e-01, -8.06326601e-01, -6.23828164e-01,
 -5.69001368e-01, -3.04082018e-01,  1.94453072e-01, -4.57292477e-01,
 -7.65464007e-01, -7.52527077e-01, -8.28836173e-01, -2.53357259e-02,
 -2.35796161e-01, -7.15492687e-01, -7.11474695e-01,  2.10213576e+00,
  1.03156652e+00,  3.18713604e-01, -8.48253604e-01, -7.62561955e-01,
  1.42656570e+00, -7.99338533e-01, -6.87549612e-01,  1.99660104e-01,
 -3.19735715e-01,  1.19280343e+00, -4.71005672e-03, -6.87025723e-01,
 -2.26540827e-01, -1.42200871e-01, -7.65832794e-01, -1.41551319e-01,
 -7.36982789e-01, -7.12046410e-01, -7.37356902e-01,  1.07778033e+00,
 -3.72941428e-01, -7.07534735e-01, -8.41389120e-01,  5.01187464e-01,
  4.34351229e-01, -6.58031923e-01, -7.65832794e-01, -5.11906997e-02,
 -3.68571479e-01, -6.97876053e-01,  9.35245520e-01,  1.49242740e+00,
  1.92834385e-01, -7.59909559e-01, -4.91101585e-01, -1.21407120e-01,
 -8.27362990e-01, -2.48095530e-01, -7.42915684e-01, -1.83714609e-03,
  5.44693957e-01, -1.37774482e-03,  7.29276013e-02,  3.84381501e-02,
 -6.67184549e-01, -2.15777927e-01, -5.07511107e-01, -7.65832794e-01,
  1.58256394e-01,  8.64780309e-01, -4.31981367e-04, -7.08398454e-01,
 -7.17891164e-01, -4.11135409e-01, -7.42280232e-01,  8.59403563e-01,
 -6.91003407e-01, -7.17330634e-01, -1.66081899e-01,  9.02431611e-01,
 -7.38390655e-01, -7.10213148e-01, -3.25870660e-01, -7.65832794e-01,
  2.29356510e+00,  3.14128139e+00, -5.57658616e-01,  2.53039597e-01,
 -7.56729517e-01,  1.25467164e+00, -7.65832794e-01, -2.08390746e-01,
  8.83837407e-02,  2.70036811e+00,  2.55891881e+00, -8.76021200e-01,
 -4.53256699e-01, -5.58174074e-01,  3.26537040e-01, -2.66865364e-01,
```

```

-1.42973593e-02, -5.83407619e-01, 7.56881625e-01, -7.08319550e-01,
-6.68939269e-01, -3.36208066e-01, -2.01978218e-01, -6.45651430e-01,
-2.18308901e-01, 1.94926574e-01, 1.80585012e+00, -3.92179169e-01,
-5.74212911e-01, -7.51610139e-01, -6.76016864e-01, -7.65832794e-01,
2.07540895e+00, -5.16695646e-01, -7.65832794e-01, 1.89508900e+00,
-7.76709486e-03, -6.98517218e-01, -6.82537022e-01, 4.37428033e-01,
-8.32533458e-01, 3.49500904e-01, -7.32391763e-01, -7.19469866e-01,
-7.11393826e-01, -7.82992886e-01, 8.53381751e-02, -3.99923381e-01,
-7.00936108e-01, 4.57452360e-01, -4.33808057e-01, 2.76150637e+00,
-1.30241440e-01, 2.08175888e-01, 3.49118218e+00, -4.60171850e-01,
-7.65115212e-01, -6.92451425e-01, -7.06119886e-01, -7.65832794e-01,
7.21219574e-01, 1.77201464e+00, -7.56418792e-01, -8.33618025e-01,
-5.51140392e-01, -7.14381782e-01, -4.58358275e-01, -3.87069288e-01,
-4.17852835e-01, 1.78655413e+00, 8.34402332e-01, -7.36581721e-01,
1.23388937e+00, -1.52320398e-01, -7.38804348e-01, -4.02751805e-01,
-7.12852354e-01, -7.65832794e-01, -5.93590916e-01, -6.58025805e-01,
-7.65832794e-01, -2.67009615e-01, 2.28632380e+00, -7.54875148e-01,
5.54432535e-01, -6.71762480e-01, 1.02544218e-01, -7.12284354e-01,
-6.69007079e-01, 1.84080047e+00, -2.80823785e-01, -7.15302063e-01,
1.26566624e-01, -8.00041631e-01, 8.97494820e-01, -8.09218514e-01,
-5.68102056e-01, 9.21034091e-01, 1.15339465e+00, -5.13886178e-01,
3.08604900e-01, -9.19145628e-01, -7.58798357e-01, -9.86060175e-01,
7.35292244e-01, -4.97683064e-02, 1.95946820e+00, -4.51691253e-01,
-3.06762760e-01, -4.17742489e-01, 8.28021569e-01, -7.10714769e-01,
1.64612455e+00, 6.51178272e-02, -2.38818664e-01, -7.70545124e-01,
2.60714368e+00, -7.65832794e-01, 1.65770345e-01, -2.44689054e-01,
-6.57996806e-01, -1.74196515e-01, 1.26083641e+00, 1.09634395e+00,
-3.65425868e-01, -1.13273447e-01, -4.26978921e-01, -4.80520224e-01,
-6.26061361e-01, 2.40505535e-01, -6.89939848e-01, 1.90847608e-01,
2.93420362e+00, 1.16339066e+00, -7.39146881e-01, -6.92185022e-01,
-6.78141635e-01, -7.83185410e-01, -7.36225339e-01, 3.11554466e-01,
-1.54010887e-01, 6.79193868e-02, -7.65832794e-01, 3.28950300e+00,
6.10945046e-01, 1.29551267e+00, -1.69389578e-01, -1.99917627e-02,
-9.43371885e-01, -7.25476738e-01, -1.63588114e-01, 2.20374461e+00,
1.10173731e+00, -3.90987156e-01, 2.31911046e+00, -4.28625756e-01,
-7.65832794e-01, -3.74989767e-01, -7.41337437e-01, -7.08158249e-01,
1.63900264e+00, -6.95840989e-01, -7.06222697e-01, -7.56319287e-01,
8.63734640e-01, 1.66371905e+00, -7.43226451e-01, -6.89579112e-01])

```

```
[208]: r3.score(X_train, y_train)
```

```
[208]: 0.8887536213334735
```

```
[209]: r3.score(X_valid, y_valid)
```

```
[209]: 0.8830729021882556
```

Higher values of 0.88 is indicative of a good NN.

Comparing the three models the best model is model 2 with a value of 0.88, followed by model 3 and then model 1. All three models have a very small difference and are generally good performing models. However Model 2 comprising of a single layer with 5 nodes performs the best as it has a value which is closest to 1.